

Q1.

Step 1: Create dimensions and fact tables.

```
CREATE TABLE customer_dim (  
    customerid integer primary key,  
    l_name character(20) NOT NULL,  
    f_name character(20),  
    city character(15) NOT NULL,  
    district character(15) NOT NULL,  
    country character(15) NOT NULL,  
    CONSTRAINT customer_customerid CHECK ((customerid > 0))  
);
```

```
CREATE Table time_dim(  
    timeid SERIAL primary key,  
    orderdate date not null,  
    dayofweek char(10) not null,  
    month char(10) not null,  
    year int not null  
);
```

```
CREATE TABLE book_dim (  
    isbn integer primary key,  
    title character(60) NOT NULL,  
    edition_no smallint DEFAULT 1,  
    price numeric(6,2) NOT NULL,  
    CONSTRAINT book_edition_no CHECK ((edition_no > 0)),  
    CONSTRAINT book_price CHECK ((price > (0)::numeric))  
);
```

```

CREATE table sales_fact(
    customerid int NOT NULL,
    timeid int NOT NULL,
    isbn int NOT NULL,
    amnt numeric(6,2) NOT NULL,
    FOREIGN KEY (customerid) REFERENCES customer_dim(customerid),
    FOREIGN KEY (timeid) REFERENCES time_dim(timeid),
    FOREIGN KEY (isbn) REFERENCES book_dim(isbn),
    CONSTRAINT amnt CHECK ((amnt > (0)::numeric))
);

ALTER TABLE ONLY sales_fact
    ADD CONSTRAINT sales_fact_pkey PRIMARY KEY (customerid, timeid, isbn);

```

Table list in DB:

```
lijins=> \dt;
```

| List of relations | | | |
|-------------------|--------------|-------|--------|
| Schema | Name | Type | Owner |
| public | author | table | lijins |
| public | book | table | lijins |
| public | book_author | table | lijins |
| public | book_dim | table | lijins |
| public | cust_order | table | lijins |
| public | customer | table | lijins |
| public | customer_dim | table | lijins |
| public | order_detail | table | lijins |
| public | sales_fact | table | lijins |
| public | time_dim | table | lijins |

(10 rows)

- book_dim is book dimension
- customer_dim is customer dimension
- time_dim is time dimension
- sales_fact is sales fact table

Step 2.Extract data from existing tables

1) customer dimension

```
INSERT INTO customer_dim (customerid, l_name, f_name, city, district, country)
( SELECT customerid,
    l_name,
    f_name,
    city,
    district,
    country
FROM customer);
```

Tijns=>INSERT 0 118

2) time dimension

```
INSERT INTO time_dim (orderdate, dayofweek, MONTH, YEAR)
( SELECT DISTINCT orderdate,
    to_char(orderdate,'Day'),
    to_char(orderdate,'Month'),
    to_number(to_char(orderdate,'YYYY'),'9999')
FROM cust_order
ORDER BY orderdate);
```

Tijns=>INSERT 0 124

3) book dimension

```
INSERT INTO book_dim (isbn, title, edition_no, price)
( SELECT isbn, title, edition_no, price FROM book);
```

Tijns=>INSERT 0 12

3) sales fact table

```
INSERT INTO sales_fact (customerid, timeid, isbn, amnt)
( SELECT cust_order_book_detail.customerid,
    time_dim.timeid,
    cust_order_book_detail.isbn,
```

```

    cust_order_book_detail.quantity * cust_order_book_detail.price AS amnt
FROM
  (SELECT cust_order_detail.customerid,
    cust_order_detail.orderdate,
    cust_order_detail.isbn,
    cust_order_detail.quantity,
    book_dim.price
  FROM
    (SELECT cust_order.customerid,
      cust_order.orderdate,
      order_detail.isbn,
      sum(order_detail.quantity) quantity
    FROM cust_order
    NATURAL JOIN order_detail
    GROUP BY cust_order.customerid,
      cust_order.orderdate,
      order_detail.isbn) AS cust_order_detail
  NATURAL JOIN book_dim) AS cust_order_book_detail
INNER JOIN time_dim ON cust_order_book_detail.orderdate = time_dim.orderdate);

```

lijins=>INSERT 0 1070

Q2.

```

create materialized view avg_amnt_view as
  select customerid, avg(amnt) as avg_amnt from sales_fact
  group by customerid;

```

lijins=>SELECT 104

Query1: select avg(avg_amnt) from avg_amnt_view;

```
lijins=> select avg(avg_amnt) from avg_amnt_v
          avg
-----
202.9588687852809865
(1 row)
```

Query2: select avg(amnt) from sales_fact;

```
lijins=> select avg(amnt) from sales_fact;
          avg
-----
161.3691588785046729
(1 row)
```

The result of query1 is the average money spent per customer.

The result of query2 is the average money spent per customer, per orderdate, per book.isbn.

So the result of query2 is correct.

Q3.

a) Using "Roll-up" and "Dice" operations.

```
SELECT customer_dim.customerid,  
       customer_dim.l_name,  
       customer_dim.f_name  
FROM customer_dim  
NATURAL JOIN  
(SELECT customerid,  
       sum(amnt) AS sum_amnt,  
       rank() over (ORDER BY sum(amnt) DESC) rank  
FROM sales_fact  
GROUP BY customerid) AS tmp_view  
WHERE tmp_view.rank < 6;
```

| customerid | l_name | f_name |
|------------|--------|--------|
| 1 | Jacson | Kirk |
| 2 | Leow | May-N |
| 3 | Andree | Peter |
| 14 | Anslow | Craig |
| 79 | Liang | Jiajun |

(5 rows)

b) Using "Slice" and "Roll-up" operation.

```
CREATE OR REPLACE FUNCTION best_perc_of_ord() RETURNS numeric AS $$  
DECLARE  
sum_amnt numeric;  
ord_count int;  
avg_ord_amnt numeric;  
per_of_ord numeric;  
greater_count int;  
no_of_ord int;  
BEGIN  
--get customer rank by amnt  
drop view if exists cust_sum_amnt_count_view;
```

```

create view cust_sum_amnt_count_view as

select customerid, sum(amnt) as sum_amnt, rank() over (order by sum(amnt) desc) rank, c
ount(1) from sales_fact

group by customerid;

--caculate amount of per order, per customer

drop view if exists cust_ord_amnt_view;

CREATE VIEW cust_ord_amnt_view as

SELECT cust_order.customerid,
       ord_book_amnt.orderid,
       SUM(ord_book_amnt.amnt) sum_amnt
FROM   cust_order
       NATURAL join (SELECT order_detail.orderid,
                           order_detail.isbn,
                           SUM(order_detail.quantity) * book_dim.price AS
                           amnt
                        FROM   order_detail
                        NATURAL join book_dim
                        GROUP BY order_detail.orderid,
                                order_detail.isbn,
                                book_dim.price) AS ord_book_amnt
GROUP BY cust_order.customerid,
         ord_book_amnt.orderid;

sum_amnt = (SELECT sum(amnt) AS sum_amnt FROM sales_fact);
ord_count = (SELECT count(1) AS ord_count FROM cust_order);
avg_ord_amnt = sum_amnt/ord_count;

--debug: raise exception 'avg_ord_amnt: %', avg_ord_amnt;

--caculate count of by the best buyer orders whose amount greater than avg_ord_amnt
greater_count = (select count(1)
from cust_ord_amnt_view inner join cust_sum_amnt_count_view
on cust_ord_amnt_view.customerid = cust_sum_amnt_count_view.customerid
where cust_sum_amnt_count_view.rank = 1
and cust_ord_amnt_view.sum_amnt > avg_ord_amnt);

-- debug: raise exception 'greater_count: %', greater_count;

```

```

-- caculate count of order by the best buyer
no_of_ord = (select count(1) from cust_order natural join cust_sum_amnt_count_view
where cust_sum_amnt_count_view.rank = 1);
--debug::raise exception 'no_of_ord: %', no_of_ord;
-- caculate count of order by best costomer
per_of_ord = greater_count/no_of_ord :: numeric;
-- debug::raise exception 'per_of_ord: %', per_of_ord;
RETURN per_of_ord;
END;
$$ LANGUAGE plpgsql;

```

```
SELECT best_perc_of_ord();
```

```

      best_perc_of_ord
-----
0.71428571428571428571
(1 row)

```

Because best_perc_of_ord is 71.4% which is between 50% and 75% , we estimate that "the best buyer has issued a greater (than average) to medium number of orders with greater (than average) amounts of money.

Q4)

-- Create view1

DROP materialized VIEW IF EXISTS View1 CASCADE;

CREATE MATERIALIZED VIEW View1 AS

SELECT c.CustomerId, F_Name, L_Name, District, TimeId,

DayOfWeek, ISBN, Amnt

FROM Sales_fact NATURAL JOIN Customer_dim c NATURAL JOIN Time_dim;

-- Create view2

DROP materialized VIEW IF EXISTS View2 CASCADE;

CREATE MATERIALIZED VIEW View2 AS

SELECT c.CustomerId, F_Name, L_Name, Year, SUM(Amnt)

FROM Sales_fact NATURAL JOIN Customer_dim c NATURAL JOIN Time_dim

GROUP BY c.CustomerId, F_Name, L_Name, Year;

-- Create view3

DROP materialized VIEW IF EXISTS View3 CASCADE;

CREATE MATERIALIZED VIEW View3 AS

SELECT District, TimeId, DayOfWeek, ISBN, SUM(Amnt)

FROM Sales_fact NATURAL JOIN Customer_dim NATURAL JOIN Time_Dim

GROUP BY District, TimeId, DayOfWeek, ISBN;

a)

--1. The "Book Orders Database"

EXPLAIN ANALYZE

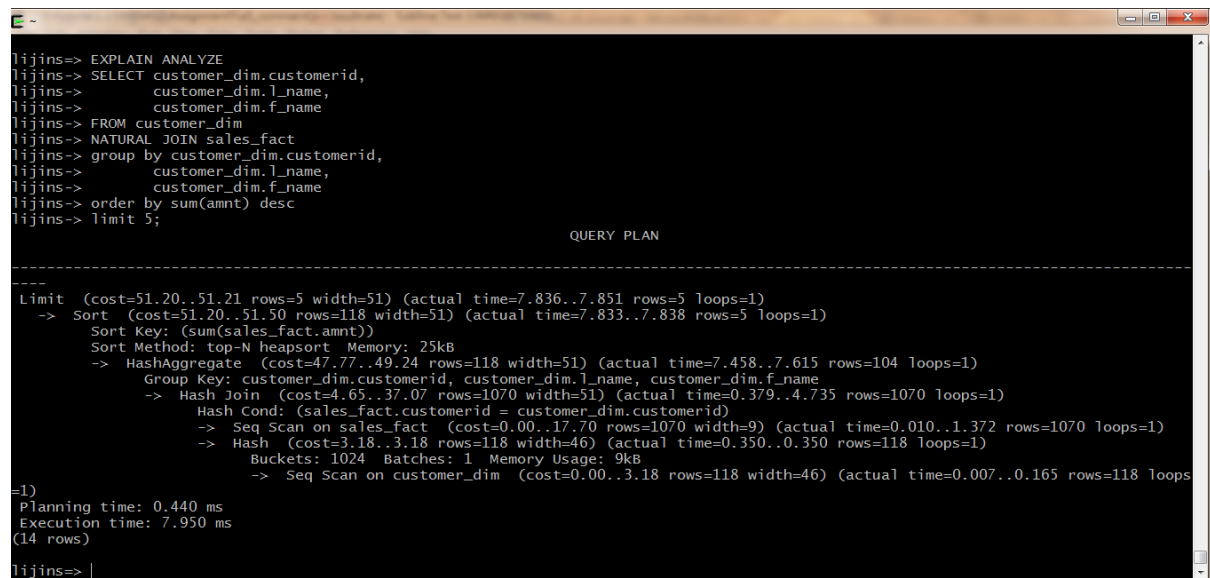
```
SELECT c.customerid,  
       sum(b.price*od.quantity)  
FROM customer c  
NATURAL JOIN cust_order co  
NATURAL JOIN order_detail od  
NATURAL JOIN book b  
GROUP BY c.customerid  
ORDER BY sum(b.price*od.quantity) DESC  
LIMIT 1;
```

```
-----  
Limit (cost=78.29..78.29 rows=1 width=20) (actual time=12.287..12.289 rows=1 loops=1)  
-> Sort (cost=78.29..78.58 rows=118 width=20) (actual time=12.282..12.282 rows=1 loops=1)  
    Sort Key: (sum((b.price * (od.quantity)::numeric)))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=76.22..77.70 rows=118 width=20) (actual time=11.940..12.092 rows=104 loops=1)  
    Group Key: c.customerid  
-> Hash Join (cost=15.97..65.22 rows=1100 width=20) (actual time=1.805..9.255 rows=1100 loops=1)  
    Hash Cond: (od.isbn = b.isbn)  
-> Hash Join (cost=14.70..48.83 rows=1100 width=10) (actual time=1.689..6.189 rows=1100 loops=1)  
    Hash Cond: (od.orderid = co.orderid)  
-> Seq Scan on order_detail od (cost=0.00..19.00 rows=1100 width=10) (actual time=0.008..1.426 rows=1100 loops=1)  
-> Hash (cost=11.93..11.93 rows=222 width=8) (actual time=1.665..1.665 rows=222 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 9kB  
-> Hash Join (cost=4.65..11.93 rows=222 width=8) (actual time=0.426..1.311 rows=222 loops=1)  
    Hash Cond: (co.customerid = c.customerid)  
-> Seq Scan on cust_order co (cost=0.00..4.22 rows=222 width=8) (actual time=0.009..0.280 rows=222 loops=1)  
-> Hash (cost=3.18..3.18 rows=118 width=4) (actual time=0.355..0.355 rows=118 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 5kB  
-> Seq Scan on customer c (cost=0.00..3.18 rows=118 width=4) (actual time=0.007..0.165 rows=118 loops=1)  
-> Hash (cost=1.12..1.12 rows=12 width=18) (actual time=0.050..0.050 rows=12 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 1kB  
-> Seq Scan on book b (cost=0.00..1.12 rows=12 width=18) (actual time=0.009..0.024 rows=12 loops=1)  
Planning time: 1.246 ms  
Execution time: 12.412 ms  
(24 rows)
```

--2. The "The Data Mart"

EXPLAIN ANALYZE

```
SELECT customer_dim.customerid,  
       customer_dim.l_name,  
       customer_dim.f_name  
FROM customer_dim  
NATURAL JOIN sales_fact  
GROUP BY customer_dim.customerid,  
         customer_dim.l_name,  
         customer_dim.f_name  
ORDER BY sum(amnt) DESC  
LIMIT 5;
```

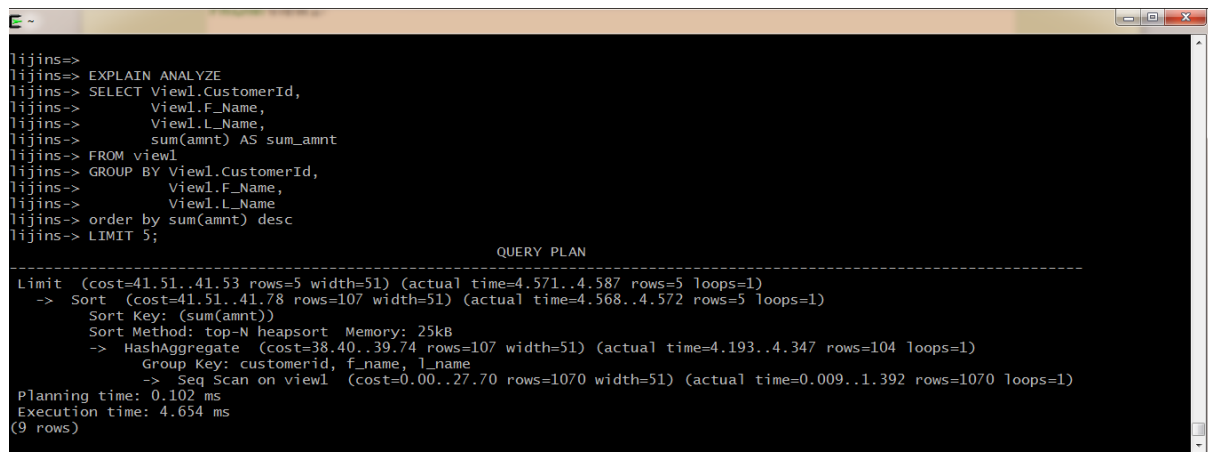


```
lijins=> EXPLAIN ANALYZE  
lijins-> SELECT customer_dim.customerid,  
lijins->       customer_dim.l_name,  
lijins->       customer_dim.f_name  
lijins-> FROM customer_dim  
lijins-> NATURAL JOIN sales_fact  
lijins-> group by customer_dim.customerid,  
lijins->       customer_dim.l_name,  
lijins->       customer_dim.f_name  
lijins-> order by sum(amnt) desc  
lijins-> limit 5;  
  
QUERY PLAN  
  
-----  
Limit (cost=51.20..51.21 rows=5 width=51) (actual time=7.836..7.851 rows=5 loops=1)  
-> Sort (cost=51.20..51.50 rows=118 width=51) (actual time=7.833..7.838 rows=5 loops=1)  
    Sort Key: (sum(sales_fact.amnt))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=47.77..49.24 rows=118 width=51) (actual time=7.458..7.615 rows=104 loops=1)  
    Group Key: customer_dim.customerid, customer_dim.l_name, customer_dim.f_name  
-> Hash Join (cost=4.65..37.07 rows=1070 width=51) (actual time=0.379..4.735 rows=1070 loops=1)  
    Hash Cond: (sales_fact.customerid = customer_dim.customerid)  
-> Seq Scan on sales_fact (cost=0.00..17.70 rows=1070 width=9) (actual time=0.010..1.372 rows=1070 loops=1)  
-> Hash (cost=3.18..3.18 rows=118 width=46) (actual time=0.350..0.350 rows=118 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 9kB  
-> Seq Scan on customer_dim (cost=0.00..3.18 rows=118 width=46) (actual time=0.007..0.165 rows=118 loops=1)  
Planning time: 0.440 ms  
Execution time: 7.950 ms  
(14 rows)  
lijins=> |
```

--3. The view View1

EXPLAIN ANALYZE

```
SELECT View1.CustomerId,  
       View1.F_Name,  
       View1.L_Name,  
       sum(amnt) AS sum_amnt  
FROM view1  
GROUP BY View1.CustomerId,  
         View1.F_Name,  
         View1.L_Name  
ORDER BY sum(amnt) DESC  
LIMIT 5;
```

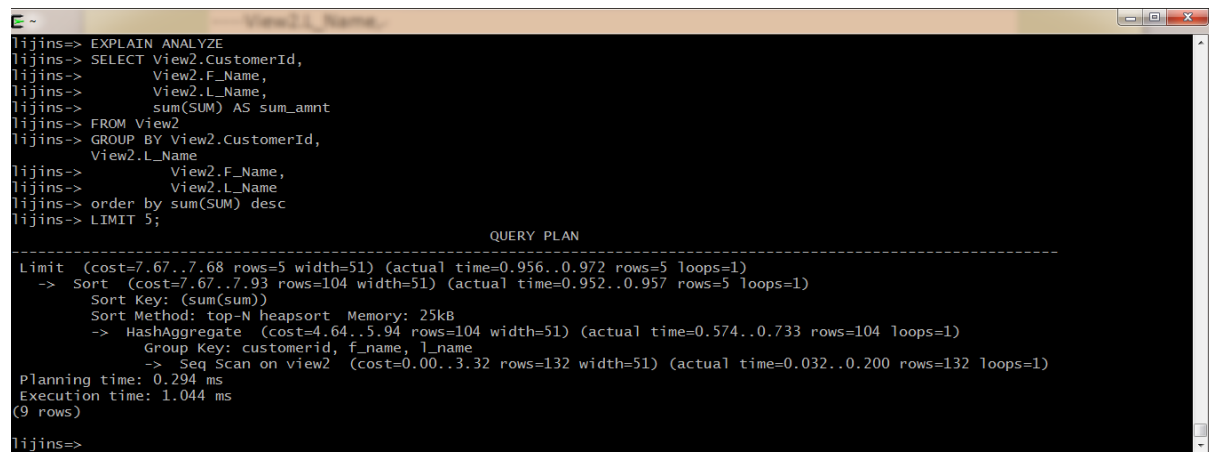


```
lijins=>  
lijins=> EXPLAIN ANALYZE  
lijins-> SELECT View1.CustomerId,  
lijins->         View1.F_Name,  
lijins->         View1.L_Name,  
lijins->         sum(amnt) AS sum_amnt  
lijins-> FROM view1  
lijins-> GROUP BY View1.CustomerId,  
lijins->         View1.F_Name,  
lijins->         View1.L_Name  
lijins-> order by sum(amnt) desc  
lijins-> LIMIT 5;  
  
-----  
QUERY PLAN  
-----  
Limit (cost=41.51..41.53 rows=5 width=51) (actual time=4.571..4.587 rows=5 loops=1)  
-> Sort (cost=41.51..41.78 rows=107 width=51) (actual time=4.568..4.572 rows=5 loops=1)  
    Sort Key: (sum(amnt))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=38.40..39.74 rows=107 width=51) (actual time=4.193..4.347 rows=104 loops=1)  
    Group Key: customerid, f_name, l_name  
-> Seq Scan on view1 (cost=0.00..27.70 rows=1070 width=51) (actual time=0.009..1.392 rows=1070 loops=1)  
Planning time: 0.102 ms  
Execution time: 4.654 ms  
(9 rows)
```

--4. The view View2

EXPLAIN ANALYZE

```
SELECT View2.CustomerId,  
       View2.F_Name,  
       View2.L_Name,  
       sum(SUM) AS sum_amnt  
FROM View2  
GROUP BY View2.CustomerId,  
         View2.F_Name,  
         View2.L_Name  
ORDER BY sum(SUM) DESC  
LIMIT 5;
```



```
lijins=> EXPLAIN ANALYZE  
lijins-> SELECT View2.CustomerId,  
lijins->         View2.F_Name,  
lijins->         View2.L_Name,  
lijins->         sum(SUM) AS sum_amnt  
lijins-> FROM View2  
lijins-> GROUP BY View2.CustomerId,  
lijins->         View2.L_Name,  
lijins->         View2.F_Name,  
lijins->         View2.L_Name  
lijins-> order by sum(SUM) desc  
lijins-> LIMIT 5;  
  
QUERY PLAN  
-----  
Limit (cost=7.67..7.68 rows=5 width=51) (actual time=0.956..0.972 rows=5 loops=1)  
-> Sort (cost=7.67..7.93 rows=104 width=51) (actual time=0.952..0.957 rows=5 loops=1)  
    Sort Key: (sum(sum))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=4.64..5.94 rows=104 width=51) (actual time=0.574..0.733 rows=104 loops=1)  
    Group Key: customerid, f_name, l_name  
-> Seq Scan on view2 (cost=0.00..3.32 rows=132 width=51) (actual time=0.032..0.200 rows=132 loops=1)  
Planning time: 0.294 ms  
Execution time: 1.044 ms  
(9 rows)  
lijins=>
```

Explain the findings:

- 1) The speed of the execution of the query "2. The Data Mata" is faster than "1. The Book Orders Database". The reason is that the sales table already contains customerid and amnt information, which reduces the join time.
- 2) The speed of the execution of the query "3.The view View1" is faster than "2. The Data Mata ". Because materialized views improve query performance by pre-calculating expensive join and aggregation operations.
- 3) The speed of the execution of the query "4.The view View2" is faster than "3. The view View1 ". Because View1 only includes join operations, whereas, View2 includes join and aggregation operations.

Q4-b)

--1. The "Book Orders Database"

EXPLAIN ANALYZE

SELECT c.country, **sum**(b.price*od.quantity)

FROM customer c

NATURAL JOIN cust_order co

NATURAL JOIN order_detail od

NATURAL JOIN book b

GROUP BY c.country

ORDER BY **sum**(b.price*od.quantity) **DESC**

LIMIT 1;

```
QUERY PLAN
-----
Limit (cost=76.34..76.35 rows=1 width=32) (actual time=12.282..12.284 rows=1 loops=1)
-> Sort (cost=76.34..76.36 rows=7 width=32) (actual time=12.277..12.277 rows=1 loops=1)
    Sort Key: (Sum((b.price * (od.quantity)::numeric)))
    Sort Method: top-N heapsort  Memory: 25kB
-> HashAggregate (cost=76.22..76.31 rows=7 width=32) (actual time=12.224..12.238 rows=7 loops=1)
    Group Key: c.country
-> Hash Join (cost=15.97..65.22 rows=1100 width=32) (actual time=1.767..9.322 rows=1100 loops=1)
    Hash Cond: (od.isbn = b.isbn)
-> Hash Join (cost=14.70..48.83 rows=1100 width=22) (actual time=1.702..6.298 rows=1100 loops=1)
    Hash Cond: (od.orderid = co.orderid)
-> Seq Scan on order_detail od (cost=0.00..19.00 rows=1100 width=10) (actual time=0.009..1.511
rows=1100 loops=1)
-> Hash (cost=11.93..11.93 rows=222 width=20) (actual time=1.665..1.665 rows=222 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 12kB
-> Hash Join (cost=4.65..11.93 rows=222 width=20) (actual time=0.423..1.326 rows=222 loops=1)
    Hash Cond: (co.customerid = c.customerid)
-> Seq Scan on cust_order co (cost=0.00..4.22 rows=222 width=8) (actual time=0.009..0.301 ro
s=222 loops=1)
-> Hash (cost=3.18..3.18 rows=118 width=20) (actual time=0.366..0.366 rows=118 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 6kB
-> Seq Scan on customer c (cost=0.00..3.18 rows=118 width=20) (actual time=0.006..0.18
rows=118 loops=1)
-> Hash (cost=1.12..1.12 rows=12 width=18) (actual time=0.051..0.051 rows=12 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 1kB
-> Seq Scan on book b (cost=0.00..1.12 rows=12 width=18) (actual time=0.009..0.029 rows=12 loops=1)
Planning time: 1.220 ms
Execution time: 12.512 ms
(24 rows)
lijins=> |
```

--2. The Data Mart,

EXPLAIN ANALYZE

SELECT country, **sum**(amnt)

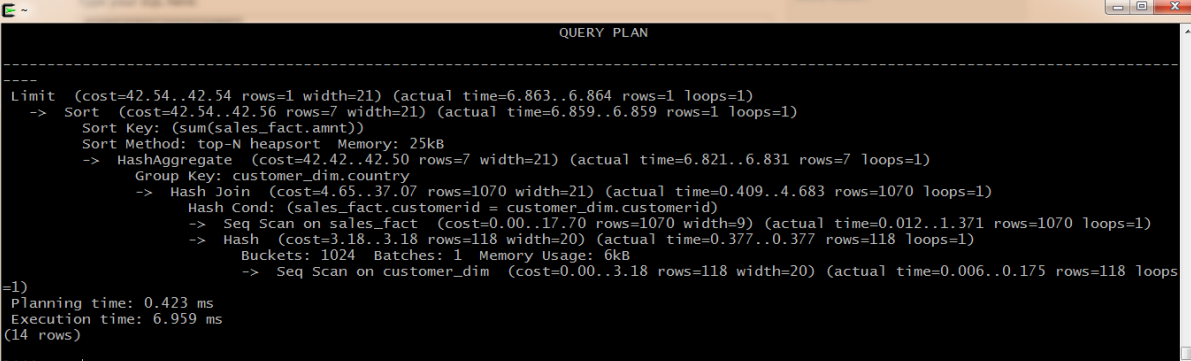
FROM customer_dim

NATURAL JOIN sales_fact

GROUP BY country

ORDER BY **sum**(amnt) **DESC**

LIMIT 1;



QUERY PLAN

```
Limit (cost=42.54..42.54 rows=1 width=21) (actual time=6.863..6.864 rows=1 loops=1)
-> Sort (cost=42.54..42.56 rows=7 width=21) (actual time=6.859..6.859 rows=1 loops=1)
    Sort Key: (sum(sales_fact.amnt))
    Sort Method: top-N heapsort  Memory: 25kB
    -> HashAggregate (cost=42.42..42.50 rows=7 width=21) (actual time=6.821..6.831 rows=7 loops=1)
        Group Key: customer_dim.country
        -> Hash Join (cost=4.65..37.07 rows=1070 width=21) (actual time=0.409..4.683 rows=1070 loops=1)
            Hash Cond: (sales_fact.customerid = customer_dim.customerid)
            -> Seq Scan on sales_fact (cost=0.00..17.70 rows=1070 width=9) (actual time=0.012..1.371 rows=1070 loops=1)
            -> Hash (cost=3.18..3.18 rows=118 width=20) (actual time=0.377..0.377 rows=118 loops=1)
                Buckets: 1024  Batches: 1  Memory Usage: 6kB
                -> Seq Scan on customer_dim (cost=0.00..3.18 rows=118 width=20) (actual time=0.006..0.175 rows=118 loops=1)
        -> Seq Scan on customer_dim (cost=0.00..3.18 rows=118 width=20) (actual time=0.006..0.175 rows=118 loops=1)
Planning time: 0.423 ms
Execution time: 6.959 ms
(14 rows)
```

lijins=> |

--3. The view View2

EXPLAIN ANALYZE

SELECT country, **sum**(sum)

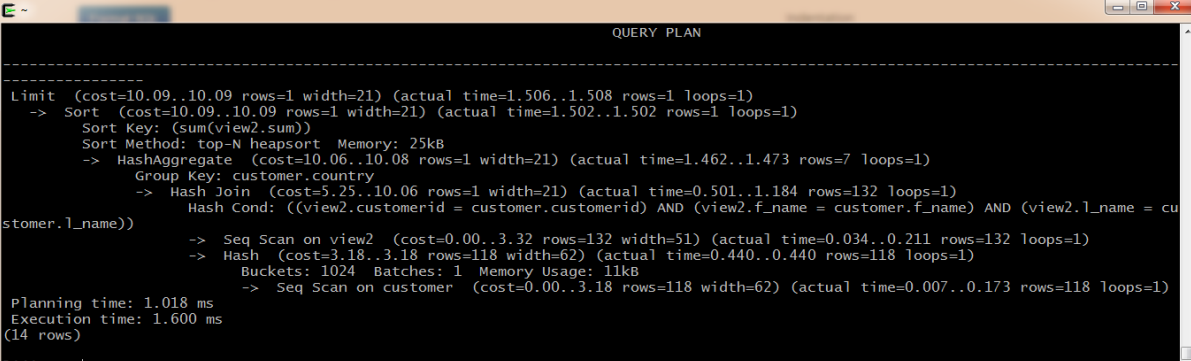
FROM View2

NATURAL JOIN customer_dim

GROUP BY country

ORDER BY **sum**(sum) **DESC**

LIMIT 1;



QUERY PLAN

```
-----
Limit (cost=10.09..10.09 rows=1 width=21) (actual time=1.506..1.508 rows=1 loops=1)
  -> Sort (cost=10.09..10.09 rows=1 width=21) (actual time=1.502..1.502 rows=1 loops=1)
        Sort Key: (sum(view2.sum))
        Sort Method: top-N heapsort  Memory: 25kB
        -> HashAggregate (cost=10.06..10.08 rows=1 width=21) (actual time=1.462..1.473 rows=7 loops=1)
              Group Key: customer.country
              -> Hash Join (cost=5.25..10.06 rows=1 width=21) (actual time=0.501..1.184 rows=132 loops=1)
                    Hash Cond: ((view2.customerid = customer.customerid) AND (view2.f_name = customer.f_name) AND (view2.l_name = customer.l_name))
                    -> Seq Scan on view2 (cost=0.00..3.32 rows=132 width=51) (actual time=0.034..0.211 rows=132 loops=1)
                          -> Hash (cost=3.18..3.18 rows=118 width=62) (actual time=0.440..0.440 rows=118 loops=1)
                                Buckets: 1024 Batches: 1 Memory Usage: 11kB
                                -> Seq Scan on customer (cost=0.00..3.18 rows=118 width=62) (actual time=0.007..0.173 rows=118 loops=1)
Planning time: 1.018 ms
Execution time: 1.600 ms
(14 rows)
lijins=> |
```


--4. The view View3.

EXPLAIN ANALYZE

```
SELECT country, sum(sum)
```

```
FROM View3
```

NATURAL JOIN

```
(SELECT DISTINCT district, country FROM customer_dim) AS tmp_cust
```

```
GROUP BY country
```

```
ORDER BY sum(sum) DESC
```

```
LIMIT 1;
```

```
-----  
Limit (cost=42.61..42.61 rows=1 width=21) (actual time=6.942..6.944 rows=1 loops=1)  
-> Sort (cost=42.61..42.65 rows=16 width=21) (actual time=6.937..6.937 rows=1 loops=1)  
    Sort Key: (sum(view3.sum))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=42.33..42.53 rows=16 width=21) (actual time=6.898..6.911 rows=7 loops=1)  
    Group Key: tmp_cust.country  
-> Hash Join (cost=4.29..37.59 rows=947 width=21) (actual time=0.539..4.935 rows=1006 loops=1)  
    Hash Cond: (view3.district = tmp_cust.district)  
-> Seq Scan on view3 (cost=0.00..20.06 rows=1006 width=21) (actual time=0.012..1.350 rows=1006 loops=1)  
-> Hash (cost=4.09..4.09 rows=16 width=32) (actual time=0.508..0.508 rows=17 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 2kB  
-> Subquery Scan on tmp_cust (cost=3.77..4.09 rows=16 width=32) (actual time=0.414..0.475 rows=17 loops=1)  
-> HashAggregate (cost=3.77..3.93 rows=16 width=32) (actual time=0.411..0.433 rows=17 loops=1)  
    Group Key: customer_dim.district, customer_dim.country  
-> Seq Scan on customer_dim (cost=0.00..3.18 rows=118 width=32) (actual time=0.009..0.171 rows=118 loops=1)  
Planning time: 0.284 ms  
Execution time: 7.045 ms  
(17 rows)  
lijins=>
```

Explain the findings:

- 1) The speed of the execution of the query "2. The Data Mata" is faster than "1. The Book Orders Database". The reason is that the sales table already contains customerid and amnt information, which reduces the join time.
- 2) The speed of the execution of the query "3.The view View2" is faster than "2. The Data Mata ". Because materialized views improve query performance by pre-calculating expensive join and aggregation operations.
- 3) The speed of the execution of the query "4.The view View3" is slower than "3. The view View2 ". We cannot use view3 directly because "district" is not a primary key, when joining to customer table, view3 needs to be modified.

Q5)-a) Display sum ammounts per customer and avg ammount per customer in a city.

```
SELECT customerid, f_name, city, sum_amnt,
       avg(sum_amnt) OVER (PARTITION BY city) AS AVG
FROM
  (SELECT customerid, f_name, city, sum(amnt) AS sum_amnt
   FROM sales_fact
   NATURAL JOIN customer_dim
   NATURAL JOIN time_dim
   WHERE (MONTH= 'April' OR MONTH= 'May') AND YEAR=2017
   GROUP BY customerid, f_name, city) AS tmp
ORDER BY city;
```

lijins-> ORDER BY city;

| customerid | f_name | city | sum_amnt | avg |
|------------|-------------|--------------|----------|---------------------------|
| 94 | Shweta | Auckland | 3615.00 | 2370.00000000000000000000 |
| 113 | Tao | Auckland | 2055.00 | 2370.00000000000000000000 |
| 95 | Priyanka | Auckland | 1440.00 | 2370.00000000000000000000 |
| 116 | Adrian | Beijing | 765.00 | 1157.50000000000000000000 |
| 107 | Xiaoxing | Beijing | 1550.00 | 1157.50000000000000000000 |
| 100 | Zoltan | Budapest | 2710.00 | 2710.00000000000000000000 |
| 99 | Nathan | Christchurch | 785.00 | 1258.33333333333333333333 |
| 98 | Valerie | Christchurch | 925.00 | 1258.33333333333333333333 |
| 115 | Christopher | Christchurch | 2065.00 | 1258.33333333333333333333 |
| 108 | Aaron | Lower Hutt | 1555.00 | 1555.00000000000000000000 |
| 118 | Daniel | Masterton | 1465.00 | 1465.00000000000000000000 |
| 101 | Benjamin | Mumbai | 1245.00 | 1245.00000000000000000000 |
| 112 | Bilal | Porirua | 1120.00 | 1297.50000000000000000000 |
| 97 | Cameron | Porirua | 1475.00 | 1297.50000000000000000000 |
| 102 | Leila | Sidney | 1165.00 | 2122.50000000000000000000 |
| 103 | Mansi | Sidney | 3080.00 | 2122.50000000000000000000 |
| 96 | Jovan | Skopje | 775.00 | 775.00000000000000000000 |
| 110 | Ronni | Upper Hutt | 1490.00 | 1490.00000000000000000000 |
| 109 | Neel | Wellington | 1440.00 | 1155.00000000000000000000 |
| 104 | Mansour | Wellington | 1425.00 | 1155.00000000000000000000 |
| 114 | Harman | Wellington | 395.00 | 1155.00000000000000000000 |
| 111 | Kaszandra | Wellington | 1360.00 | 1155.00000000000000000000 |
| 105 | Jessie | Wuhan | 1670.00 | 1776.66666666666666666667 |
| 117 | Lei | Wuhan | 1625.00 | 1776.66666666666666666667 |
| 106 | Li | Wuhan | 2035.00 | 1776.66666666666666666667 |

(25 rows)

lijins=>

Q5-b) Display daily sum and cumulative sum by orderdate in a city.

```

SELECT DISTINCT city, timeid, orderdate AS day, daily_sum, cumulative_sum
FROM
  (SELECT city, timeid, orderdate,
    sum(amnt) OVER w1 AS daily_sum,
    sum(amnt) OVER w2 AS cumulative_sum
  FROM sales_fact
  NATURAL JOIN time_dim
  NATURAL JOIN customer_dim
  WHERE (MONTH = 'April' OR MONTH = 'May') AND YEAR=2017
    WINDOW w1 AS (PARTITION BY city, orderdate),
           w2 AS (PARTITION BY city ORDER BY orderdate) ) AS tmp
ORDER BY city, orderdate;

```

| city | timeid | day | daily_sum | cumulative_sum |
|--------------|--------|------------|-----------|----------------|
| Auckland | 119 | 2017-04-23 | 360.00 | 360.00 |
| Auckland | 120 | 2017-04-29 | 2250.00 | 2610.00 |
| Auckland | 121 | 2017-04-30 | 2805.00 | 5415.00 |
| Auckland | 122 | 2017-05-05 | 1695.00 | 7110.00 |
| Beijing | 117 | 2017-04-21 | 1550.00 | 1550.00 |
| Beijing | 124 | 2017-05-15 | 765.00 | 2315.00 |
| Budapest | 111 | 2017-04-15 | 2710.00 | 2710.00 |
| Christchurch | 110 | 2017-04-14 | 925.00 | 925.00 |
| Christchurch | 111 | 2017-04-15 | 785.00 | 1710.00 |
| Christchurch | 122 | 2017-05-05 | 175.00 | 1885.00 |
| Christchurch | 123 | 2017-05-06 | 1810.00 | 3695.00 |
| Christchurch | 124 | 2017-05-15 | 80.00 | 3775.00 |
| Lower Hutt | 118 | 2017-04-22 | 1555.00 | 1555.00 |
| Masterton | 124 | 2017-05-15 | 1465.00 | 1465.00 |
| Mumbai | 112 | 2017-04-16 | 1245.00 | 1245.00 |
| Porirua | 108 | 2017-04-12 | 170.00 | 170.00 |
| Porirua | 109 | 2017-04-13 | 1135.00 | 1305.00 |
| Porirua | 110 | 2017-04-14 | 170.00 | 1475.00 |
| Porirua | 122 | 2017-05-05 | 1120.00 | 2595.00 |
| Sidney | 112 | 2017-04-16 | 850.00 | 850.00 |
| Sidney | 113 | 2017-04-17 | 1305.00 | 2155.00 |
| Sidney | 114 | 2017-04-18 | 2090.00 | 4245.00 |
| Skopje | 108 | 2017-04-12 | 775.00 | 775.00 |
| Upper Hutt | 118 | 2017-04-22 | 1490.00 | 1490.00 |
| Wellington | 114 | 2017-04-18 | 1425.00 | 1425.00 |
| Wellington | 118 | 2017-04-22 | 1440.00 | 2865.00 |
| Wellington | 119 | 2017-04-23 | 1360.00 | 4225.00 |
| Wellington | 123 | 2017-05-06 | 395.00 | 4620.00 |
| Wuhan | 115 | 2017-04-19 | 1735.00 | 1735.00 |
| Wuhan | 116 | 2017-04-20 | 1525.00 | 3260.00 |
| Wuhan | 117 | 2017-04-21 | 195.00 | 3455.00 |
| Wuhan | 118 | 2017-04-22 | 250.00 | 3705.00 |
| Wuhan | 124 | 2017-05-15 | 1625.00 | 5330.00 |

(33 rows)